

Fin Time Series Homework 2

Zhenzhen Zhang(198000379)

Question 1 a) ¶

```
In [1]: library(MASS)
        data(Boston)
```

```
In [61]: dim(Boston)
```

506 14

```
In [2]: summary(Boston)
```

crim	zn	indus	chas
Min. : 0.00632	Min. : 0.00	Min. : 0.46	Min. : 0.00000
1st Qu.: 0.08204	1st Qu.: 0.00	1st Qu.: 5.19	1st Qu.: 0.00000
Median : 0.25651	Median : 0.00	Median : 9.69	Median : 0.00000
Mean : 3.61352	Mean : 11.36	Mean : 11.14	Mean : 0.06917
3rd Qu.: 3.67708	3rd Qu.: 12.50	3rd Qu.: 18.10	3rd Qu.: 0.00000
Max. : 88.97620	Max. : 100.00	Max. : 27.74	Max. : 1.00000

nox	rm	age	dis
Min. : 0.3850	Min. : 3.561	Min. : 2.90	Min. : 1.130
1st Qu.: 0.4490	1st Qu.: 5.886	1st Qu.: 45.02	1st Qu.: 2.100
Median : 0.5380	Median : 6.208	Median : 77.50	Median : 3.207
Mean : 0.5547	Mean : 6.285	Mean : 68.57	Mean : 3.795
3rd Qu.: 0.6240	3rd Qu.: 6.623	3rd Qu.: 94.08	3rd Qu.: 5.188
Max. : 0.8710	Max. : 8.780	Max. : 100.00	Max. : 12.127

rad	tax	ptratio	black
Min. : 1.000	Min. : 187.0	Min. : 12.60	Min. : 0.32
1st Qu.: 4.000	1st Qu.: 279.0	1st Qu.: 17.40	1st Qu.: 375.38
Median : 5.000	Median : 330.0	Median : 19.05	Median : 391.44
Mean : 9.549	Mean : 408.2	Mean : 18.46	Mean : 356.67
3rd Qu.: 24.000	3rd Qu.: 666.0	3rd Qu.: 20.20	3rd Qu.: 396.23
Max. : 24.000	Max. : 711.0	Max. : 22.00	Max. : 396.90

lstat	medv
Min. : 1.73	Min. : 5.00
1st Qu.: 6.95	1st Qu.: 17.02
Median : 11.36	Median : 21.20
Mean : 12.65	Mean : 22.53
3rd Qu.: 16.95	3rd Qu.: 25.00
Max. : 37.97	Max. : 50.00

```
In [3]: lmod<-lm(medv~crim+zn+indus+chas+nox+rm+age+dis+rad+tax+ptratio+black+lstat,da
        ta=Boston)
```

In [4]: `summary(lmod)`

Call:

```
lm(formula = medv ~ crim + zn + indus + chas + nox + rm + age +
    dis + rad + tax + ptratio + black + lstat, data = Boston)
```

Residuals:

Min	1Q	Median	3Q	Max
-15.595	-2.730	-0.518	1.777	26.199

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.646e+01	5.103e+00	7.144	3.28e-12	***
crim	-1.080e-01	3.286e-02	-3.287	0.001087	**
zn	4.642e-02	1.373e-02	3.382	0.000778	***
indus	2.056e-02	6.150e-02	0.334	0.738288	
chas	2.687e+00	8.616e-01	3.118	0.001925	**
nox	-1.777e+01	3.820e+00	-4.651	4.25e-06	***
rm	3.810e+00	4.179e-01	9.116	< 2e-16	***
age	6.922e-04	1.321e-02	0.052	0.958229	
dis	-1.476e+00	1.995e-01	-7.398	6.01e-13	***
rad	3.060e-01	6.635e-02	4.613	5.07e-06	***
tax	-1.233e-02	3.760e-03	-3.280	0.001112	**
ptratio	-9.527e-01	1.308e-01	-7.283	1.31e-12	***
black	9.312e-03	2.686e-03	3.467	0.000573	***
lstat	-5.248e-01	5.072e-02	-10.347	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.745 on 492 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7338

F-statistic: 108.1 on 13 and 492 DF, p-value: < 2.2e-16

From the summary of the multiple regression, we can see that zn,nox,rm,dis,rad,ptratio,black,lstat are statistically significant given a 99.9% confidence interval. Similarly, crim, chas, tax are significant given a 99% confidence interval. For indus and age, they are not statistically significant.

```
In [21]: step1<-step(lmod,direction="backward")
```

Start: AIC=1589.64

medv ~ crim + zn + indus + chas + nox + rm + age + dis + rad +
tax + ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- age	1	0.06	11079	1587.7
- indus	1	2.52	11081	1587.8
<none>			11079	1589.6
- chas	1	218.97	11298	1597.5
- tax	1	242.26	11321	1598.6
- crim	1	243.22	11322	1598.6
- zn	1	257.49	11336	1599.3
- black	1	270.63	11349	1599.8
- rad	1	479.15	11558	1609.1
- nox	1	487.16	11566	1609.4
- ptratio	1	1194.23	12273	1639.4
- dis	1	1232.41	12311	1641.0
- rm	1	1871.32	12950	1666.6
- lstat	1	2410.84	13490	1687.3

Step: AIC=1587.65

medv ~ crim + zn + indus + chas + nox + rm + dis + rad + tax +
ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- indus	1	2.52	11081	1585.8
<none>			11079	1587.7
- chas	1	219.91	11299	1595.6
- tax	1	242.24	11321	1596.6
- crim	1	243.20	11322	1596.6
- zn	1	260.32	11339	1597.4
- black	1	272.26	11351	1597.9
- rad	1	481.09	11560	1607.2
- nox	1	520.87	11600	1608.9
- ptratio	1	1200.23	12279	1637.7
- dis	1	1352.26	12431	1643.9
- rm	1	1959.55	13038	1668.0
- lstat	1	2718.88	13798	1696.7

Step: AIC=1585.76

medv ~ crim + zn + chas + nox + rm + dis + rad + tax + ptratio +
black + lstat

	Df	Sum of Sq	RSS	AIC
<none>			11081	1585.8
- chas	1	227.21	11309	1594.0
- crim	1	245.37	11327	1594.8
- zn	1	257.82	11339	1595.4
- black	1	270.82	11352	1596.0
- tax	1	273.62	11355	1596.1
- rad	1	500.92	11582	1606.1
- nox	1	541.91	11623	1607.9
- ptratio	1	1206.45	12288	1636.0
- dis	1	1448.94	12530	1645.9
- rm	1	1963.66	13045	1666.3
- lstat	1	2723.48	13805	1695.0

```
In [22]: summary(step1)
```

Call:

```
lm(formula = medv ~ crim + zn + chas + nox + rm + dis + rad +  
    tax + ptratio + black + lstat, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.5984	-2.7386	-0.5046	1.7273	26.2373

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	36.341145	5.067492	7.171	2.73e-12	***
crim	-0.108413	0.032779	-3.307	0.001010	**
zn	0.045845	0.013523	3.390	0.000754	***
chas	2.718716	0.854240	3.183	0.001551	**
nox	-17.376023	3.535243	-4.915	1.21e-06	***
rm	3.801579	0.406316	9.356	< 2e-16	***
dis	-1.492711	0.185731	-8.037	6.84e-15	***
rad	0.299608	0.063402	4.726	3.00e-06	***
tax	-0.011778	0.003372	-3.493	0.000521	***
ptratio	-0.946525	0.129066	-7.334	9.24e-13	***
black	0.009291	0.002674	3.475	0.000557	***
lstat	-0.522553	0.047424	-11.019	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.736 on 494 degrees of freedom

Multiple R-squared: 0.7406, Adjusted R-squared: 0.7348

F-statistic: 128.2 on 11 and 494 DF, p-value: < 2.2e-16

```
In [23]: step2<-step(lmod,direction="backward", k=log(nrow(Boston)))
```

Start: AIC=1648.81

medv ~ crim + zn + indus + chas + nox + rm + age + dis + rad +
tax + ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- age	1	0.06	11079	1642.6
- indus	1	2.52	11081	1642.7
<none>			11079	1648.8
- chas	1	218.97	11298	1652.5
- tax	1	242.26	11321	1653.5
- crim	1	243.22	11322	1653.6
- zn	1	257.49	11336	1654.2
- black	1	270.63	11349	1654.8
- rad	1	479.15	11558	1664.0
- nox	1	487.16	11566	1664.4
- ptratio	1	1194.23	12273	1694.4
- dis	1	1232.41	12311	1696.0
- rm	1	1871.32	12950	1721.6
- lstat	1	2410.84	13490	1742.2

Step: AIC=1642.59

medv ~ crim + zn + indus + chas + nox + rm + dis + rad + tax +
ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- indus	1	2.52	11081	1636.5
<none>			11079	1642.6
- chas	1	219.91	11299	1646.3
- tax	1	242.24	11321	1647.3
- crim	1	243.20	11322	1647.3
- zn	1	260.32	11339	1648.1
- black	1	272.26	11351	1648.7
- rad	1	481.09	11560	1657.9
- nox	1	520.87	11600	1659.6
- ptratio	1	1200.23	12279	1688.4
- dis	1	1352.26	12431	1694.6
- rm	1	1959.55	13038	1718.8
- lstat	1	2718.88	13798	1747.4

Step: AIC=1636.48

medv ~ crim + zn + chas + nox + rm + dis + rad + tax + ptratio +
black + lstat

	Df	Sum of Sq	RSS	AIC
<none>			11081	1636.5
- chas	1	227.21	11309	1640.5
- crim	1	245.37	11327	1641.3
- zn	1	257.82	11339	1641.9
- black	1	270.82	11352	1642.5
- tax	1	273.62	11355	1642.6
- rad	1	500.92	11582	1652.6
- nox	1	541.91	11623	1654.4
- ptratio	1	1206.45	12288	1682.5
- dis	1	1448.94	12530	1692.4
- rm	1	1963.66	13045	1712.8
- lstat	1	2723.48	13805	1741.5

```
In [26]: summary(step2)
```

Call:

```
lm(formula = medv ~ crim + zn + chas + nox + rm + dis + rad +  
    tax + ptratio + black + lstat, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-15.5984	-2.7386	-0.5046	1.7273	26.2373

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	36.341145	5.067492	7.171	2.73e-12	***
crim	-0.108413	0.032779	-3.307	0.001010	**
zn	0.045845	0.013523	3.390	0.000754	***
chas	2.718716	0.854240	3.183	0.001551	**
nox	-17.376023	3.535243	-4.915	1.21e-06	***
rm	3.801579	0.406316	9.356	< 2e-16	***
dis	-1.492711	0.185731	-8.037	6.84e-15	***
rad	0.299608	0.063402	4.726	3.00e-06	***
tax	-0.011778	0.003372	-3.493	0.000521	***
ptratio	-0.946525	0.129066	-7.334	9.24e-13	***
black	0.009291	0.002674	3.475	0.000557	***
lstat	-0.522553	0.047424	-11.019	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.736 on 494 degrees of freedom


Multiple R-squared: 0.7406, Adjusted R-squared: 0.7348

F-statistic: 128.2 on 11 and 494 DF, p-value: < 2.2e-16

(2) From the AIC and BIC, we can see that both AIC and BIC set indus and age aside the model and the remaining variables are all significant.


Question 1 b)


```
In [14]: hist(Boston$medv)
```



From the histogram we can know that medv is right-skewed because the distribution's peak is off center toward the limit and a tail stretches away from it. In this case, we need to modify the outcome variable. For example, $\log(\text{medv})$:

```
In [58]: hist(log(Boston$medv))
```



```
In [18]: lmod2<-lm(log(medv)~crim+zn+indus+chas+nox+rm+age+dis+rad+tax+prratio+black+ls  
tat,data=Boston)
```

```
In [20]: summary(lmod2)
```

Call:

```
lm(formula = log(medv) ~ crim + zn + indus + chas + nox + rm +
    age + dis + rad + tax + ptratio + black + lstat, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.73361	-0.09747	-0.01657	0.09629	0.86435

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	4.1020423	0.2042726	20.081	< 2e-16	***
crim	-0.0102715	0.0013155	-7.808	3.52e-14	***
zn	0.0011725	0.0005495	2.134	0.033349	*
indus	0.0024668	0.0024614	1.002	0.316755	
chas	0.1008876	0.0344859	2.925	0.003598	**
nox	-0.7783993	0.1528902	-5.091	5.07e-07	***
rm	0.0908331	0.0167280	5.430	8.87e-08	***
age	0.0002106	0.0005287	0.398	0.690567	
dis	-0.0490873	0.0079834	-6.149	1.62e-09	***
rad	0.0142673	0.0026556	5.373	1.20e-07	***
tax	-0.0006258	0.0001505	-4.157	3.80e-05	***
ptratio	-0.0382715	0.0052365	-7.309	1.10e-12	***
black	0.0004136	0.0001075	3.847	0.000135	***
lstat	-0.0290355	0.0020299	-14.304	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1899 on 492 degrees of freedom

Multiple R-squared: 0.7896, Adjusted R-squared: 0.7841

F-statistic: 142.1 on 13 and 492 DF, p-value: < 2.2e-16

```
In [27]: step3<-step(lmod2,direction="backward")
```

Start: AIC=-1667.19

log(medv) ~ crim + zn + indus + chas + nox + rm + age + dis +
rad + tax + ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- age	1	0.0057	17.755	-1669.0
- indus	1	0.0362	17.786	-1668.2
<none>			17.749	-1667.2
- zn	1	0.1643	17.914	-1664.5
- chas	1	0.3088	18.058	-1660.5
- black	1	0.5339	18.283	-1654.2
- tax	1	0.6235	18.373	-1651.7
- nox	1	0.9351	18.684	-1643.2
- rad	1	1.0413	18.791	-1640.3
- rm	1	1.0637	18.813	-1639.7
- dis	1	1.3639	19.113	-1631.7
- ptratio	1	1.9270	19.676	-1617.0
- crim	1	2.1995	19.949	-1610.1
- lstat	1	7.3809	25.130	-1493.2

Step: AIC=-1669.03

log(medv) ~ crim + zn + indus + chas + nox + rm + dis + rad +
tax + ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- indus	1	0.0363	17.791	-1670.0
<none>			17.755	-1669.0
- zn	1	0.1593	17.914	-1666.5
- chas	1	0.3138	18.069	-1662.2
- black	1	0.5431	18.298	-1655.8
- tax	1	0.6205	18.376	-1653.7
- nox	1	0.9645	18.720	-1644.3
- rad	1	1.0356	18.791	-1642.3
- rm	1	1.1452	18.900	-1639.4
- dis	1	1.5471	19.302	-1628.8
- ptratio	1	1.9224	19.677	-1619.0
- crim	1	2.1988	19.954	-1612.0
- lstat	1	8.1949	25.950	-1479.0

Step: AIC=-1670

log(medv) ~ crim + zn + chas + nox + rm + dis + rad + tax + ptratio +
black + lstat

	Df	Sum of Sq	RSS	AIC
<none>			17.791	-1670.0
- zn	1	0.1451	17.936	-1667.9
- chas	1	0.3399	18.131	-1662.4
- black	1	0.5344	18.326	-1657.0
- tax	1	0.6139	18.405	-1654.8
- nox	1	0.9350	18.726	-1646.1
- rad	1	1.0088	18.800	-1644.1
- rm	1	1.1171	18.909	-1641.2
- dis	1	1.7385	19.530	-1624.8
- ptratio	1	1.8862	19.678	-1621.0
- crim	1	2.2229	20.014	-1612.4
- lstat	1	8.1604	25.952	-1481.0

In [28]: `summary(step3)`

Call:

```
lm(formula = log(medv) ~ crim + zn + chas + nox + rm + dis +
    rad + tax + ptratio + black + lstat, data = Boston)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.73400	-0.09460	-0.01771	0.09782	0.86290

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	4.0836823	0.2030491	20.112	< 2e-16	***
crim	-0.0103187	0.0013134	-7.856	2.49e-14	***
zn	0.0010874	0.0005418	2.007	0.045308	*
chas	0.1051484	0.0342285	3.072	0.002244	**
nox	-0.7217440	0.1416535	-5.095	4.97e-07	***
rm	0.0906728	0.0162807	5.569	4.20e-08	***
dis	-0.0517059	0.0074420	-6.948	1.18e-11	***
rad	0.0134457	0.0025405	5.293	1.82e-07	***
tax	-0.0005579	0.0001351	-4.129	4.28e-05	***
ptratio	-0.0374259	0.0051715	-7.237	1.77e-12	***
black	0.0004127	0.0001071	3.852	0.000133	***
lstat	-0.0286039	0.0019002	-15.053	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1898 on 494 degrees of freedom

Multiple R-squared: 0.7891, Adjusted R-squared: 0.7844

F-statistic: 168.1 on 11 and 494 DF, p-value: < 2.2e-16

```
In [29]: step4<-step(lmod2,direction="backward", k=log(nrow(Boston)))
```

Start: AIC=-1608.02

log(medv) ~ crim + zn + indus + chas + nox + rm + age + dis +
rad + tax + ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- age	1	0.0057	17.755	-1614.1
- indus	1	0.0362	17.786	-1613.2
- zn	1	0.1643	17.914	-1609.6
<none>			17.749	-1608.0
- chas	1	0.3088	18.058	-1605.5
- black	1	0.5339	18.283	-1599.2
- tax	1	0.6235	18.373	-1596.8
- nox	1	0.9351	18.684	-1588.3
- rad	1	1.0413	18.791	-1585.4
- rm	1	1.0637	18.813	-1584.8
- dis	1	1.3639	19.113	-1576.8
- ptratio	1	1.9270	19.676	-1562.1
- crim	1	2.1995	19.949	-1555.1
- lstat	1	7.3809	25.130	-1438.3

Step: AIC=-1614.09

log(medv) ~ crim + zn + indus + chas + nox + rm + dis + rad +
tax + ptratio + black + lstat

	Df	Sum of Sq	RSS	AIC
- indus	1	0.0363	17.791	-1619.3
- zn	1	0.1593	17.914	-1615.8
<none>			17.755	-1614.1
- chas	1	0.3138	18.069	-1611.5
- black	1	0.5431	18.298	-1605.1
- tax	1	0.6205	18.376	-1602.9
- nox	1	0.9645	18.720	-1593.5
- rad	1	1.0356	18.791	-1591.6
- rm	1	1.1452	18.900	-1588.7
- dis	1	1.5471	19.302	-1578.0
- ptratio	1	1.9224	19.677	-1568.3
- crim	1	2.1988	19.954	-1561.2
- lstat	1	8.1949	25.950	-1428.3

Step: AIC=-1619.28

log(medv) ~ crim + zn + chas + nox + rm + dis + rad + tax + ptratio +
black + lstat

	Df	Sum of Sq	RSS	AIC
- zn	1	0.1451	17.936	-1621.4
<none>			17.791	-1619.3
- chas	1	0.3399	18.131	-1615.9
- black	1	0.5344	18.326	-1610.5
- tax	1	0.6139	18.405	-1608.3
- nox	1	0.9350	18.726	-1599.6
- rad	1	1.0088	18.800	-1597.6
- rm	1	1.1171	18.909	-1594.7
- dis	1	1.7385	19.530	-1578.3
- ptratio	1	1.8862	19.678	-1574.5
- crim	1	2.2229	20.014	-1565.9
- lstat	1	8.1604	25.952	-1434.5

Step: AIC=-1621.4

log(medv) ~ crim + chas + nox + rm + dis + rad + tax + ptratio +
black + lstat

	Df	Sum of Sq	RSS	AIC
<none>			17.936	-1621.4
- chas	1	0.3388	18.275	-1618.2
- tax	1	0.5229	18.459	-1613.1
- black	1	0.5386	18.475	-1612.7
- rad	1	0.9601	18.897	-1601.2
- nox	1	1.0250	18.961	-1599.5
- rm	1	1.2650	19.201	-1593.1
- dis	1	1.6967	19.633	-1581.9
- crim	1	2.1377	20.074	-1570.7
- ptratio	1	2.5632	20.500	-1560.0
- lstat	1	8.1516	26.088	-1438.1

In [30]: `summary(step4)`

Call:

`lm(formula = log(medv) ~ crim + chas + nox + rm + dis + rad +
tax + ptratio + black + lstat, data = Boston)`

Residuals:

	Min	1Q	Median	3Q	Max
	-0.73091	-0.10531	-0.01614	0.09871	0.87074

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.1000969	0.2035038	20.148	< 2e-16 ***
crim	-0.0100763	0.0013119	-7.681	8.52e-14 ***
chas	0.1049848	0.0343329	3.058	0.002350 **
nox	-0.7515379	0.1413035	-5.319	1.59e-07 ***
rm	0.0954516	0.0161547	5.909	6.43e-09 ***
dis	-0.0442395	0.0064650	-6.843	2.30e-11 ***
rad	0.0130841	0.0025418	5.148	3.81e-07 ***
tax	-0.0005050	0.0001329	-3.799	0.000164 ***
ptratio	-0.0409840	0.0048729	-8.411	4.39e-16 ***
black	0.0004143	0.0001075	3.855	0.000131 ***
lstat	-0.0285881	0.0019060	-14.999	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1904 on 495 degrees of freedom

Multiple R-squared: 0.7874, Adjusted R-squared: 0.7831

F-statistic: 183.4 on 10 and 495 DF, p-value: < 2.2e-16

From step1 to step4, these all don't use indus and age in models. And there are some obvious differences between the outcomes. To be specific, for the remaining variables: step1 in AIC: $\text{medv} \sim \text{crim} + \text{zn} + \text{chas} + \text{nox} + \text{rm} + \text{dis} + \text{rad} + \text{tax} + \text{ptratio} + \text{black} + \text{lstat}$ step2 in BIC: $\text{medv} \sim \text{crim} + \text{zn} + \text{chas} + \text{nox} + \text{rm} + \text{dis} + \text{rad} + \text{tax} + \text{ptratio} + \text{black} + \text{lstat}$ step3 in AIC: $\log(\text{medv}) \sim \text{crim} + \text{zn} + \text{chas} + \text{nox} + \text{rm} + \text{dis} + \text{rad} + \text{tax} + \text{ptratio} + \text{black} + \text{lstat}$ step4 in BIC: $\log(\text{medv}) \sim \text{crim} + \text{chas} + \text{nox} + \text{rm} + \text{dis} + \text{rad} + \text{tax} + \text{ptratio} + \text{black} + \text{lstat}$ Only step4 removes zn in model.

For adjusted R2: step1 in AIC: 0.7348 step2 in BIC: 0.7348 step3 in AIC: 0.7844 step4 in BIC: 0.7831 From above, we can know that step3 can explain more in outcome so it's better than other three.

Question 2 a)

```
In [81]: install.packages("fpp")
library("fpp")
data(fancy)
```

Installing package into 'C:/Users/zzzha/Documents/R/win-library/3.6'
(as 'lib' is unspecified)

Warning message:
"package 'fpp' is in use and will not be installed"
Warning message in data(fancy):
"data set 'fancy' not found"

```
In [38]: summary(fancy)
```


```
Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
1665   5884   8772   14316   16889  104661
```

```
In [60]: fancy
```

A Time Series: 7 × 12

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep
1987	1664.81	2397.53	2840.71	3547.29	3752.96	3714.74	4349.61	3566.34	5021.82
1988	2499.81	5198.24	7225.14	4806.03	5900.88	4951.34	6179.12	4752.15	5496.43
1989	4717.02	5702.63	9957.58	5304.78	6492.43	6630.80	7349.62	8176.62	8573.17
1990	5921.10	5814.58	12421.25	6369.77	7609.12	7224.75	8121.22	7979.25	8093.06
1991	4826.64	6470.23	9638.77	8821.17	8722.37	10209.48	11276.55	12552.22	11637.39
1992	7615.03	9849.69	14558.40	11587.33	9332.56	13082.09	16732.78	19888.61	23933.38
1993	10243.24	11266.88	21826.84	17357.33	15997.79	18601.53	26155.15	28586.52	30505.41

```
In [40]: plot.ts(fancy)
```



From the summary(fancy) and the plot, we can know that there is a increase in monthly sales every March and December. Also, the increase is gradually growing while 1991 is an exception.

Question 2 b)

```
In [87]: y <- log(fancy)
```

```
In [88]: fit <- tslm(y ~ trend + season,data=y)
```

In [89]: `summary(fit)`

Call:

`tslm(formula = y ~ trend + season, data = y)`

Residuals:

	Min	1Q	Median	3Q	Max
	-0.41644	-0.12619	0.00608	0.11389	0.38567

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	7.6058604	0.0768740	98.939	< 2e-16	***
trend	0.0223930	0.0008448	26.508	< 2e-16	***
season2	0.2510437	0.0993278	2.527	0.013718	*
season3	0.6952066	0.0993386	6.998	1.18e-09	***
season4	0.3829341	0.0993565	3.854	0.000252	***
season5	0.4079944	0.0993817	4.105	0.000106	***
season6	0.4469625	0.0994140	4.496	2.63e-05	***
season7	0.6082156	0.0994534	6.116	4.69e-08	***
season8	0.5853524	0.0995001	5.883	1.21e-07	***
season9	0.6663446	0.0995538	6.693	4.27e-09	***
season10	0.7440336	0.0996148	7.469	1.61e-10	***
season11	1.2030164	0.0996828	12.068	< 2e-16	***
season12	1.9581366	0.0997579	19.629	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1858 on 71 degrees of freedom

Multiple R-squared: 0.9527, Adjusted R-squared: 0.9447

F-statistic: 119.1 on 12 and 71 DF, p-value: < 2.2e-16

Question 2 c)

In [82]:

```

x1=1:84
x2=rep(c(0,1,0,0,0,0,0,0,0,0,0,0),7)
x3=rep(c(0,0,1,0,0,0,0,0,0,0,0,0),7)
x4=rep(c(0,0,0,1,0,0,0,0,0,0,0,0),7)
x5=rep(c(0,0,0,0,1,0,0,0,0,0,0,0),7)
x6=rep(c(0,0,0,0,0,1,0,0,0,0,0,0),7)
x7=rep(c(0,0,0,0,0,0,1,0,0,0,0,0),7)
x8=rep(c(0,0,0,0,0,0,0,1,0,0,0,0),7)
x9=rep(c(0,0,0,0,0,0,0,0,1,0,0,0),7)
x10=rep(c(0,0,0,0,0,0,0,0,0,1,0,0),7)
x11=rep(c(0,0,0,0,0,0,0,0,0,0,1,0),7)
x12=rep(c(0,0,0,0,0,0,0,0,0,0,0,1),7)

```

In [85]: `fit1<-lm(log(fancy)~x1+x2+x3+x4+x5+x6+x7+x8+x9+x10+x11+x12)`

```
In [86]: summary(fit1)
```

Call:

```
lm(formula = log(fancy) ~ x1 + x2 + x3 + x4 + x5 + x6 + x7 +
    x8 + x9 + x10 + x11 + x12)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.41644	-0.12619	0.00608	0.11389	0.38567

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	7.6058604	0.0768740	98.939	< 2e-16	***
x1	0.0223930	0.0008448	26.508	< 2e-16	***
x2	0.2510437	0.0993278	2.527	0.013718	*
x3	0.6952066	0.0993386	6.998	1.18e-09	***
x4	0.3829341	0.0993565	3.854	0.000252	***
x5	0.4079944	0.0993817	4.105	0.000106	***
x6	0.4469625	0.0994140	4.496	2.63e-05	***
x7	0.6082156	0.0994534	6.116	4.69e-08	***
x8	0.5853524	0.0995001	5.883	1.21e-07	***
x9	0.6663446	0.0995538	6.693	4.27e-09	***
x10	0.7440336	0.0996148	7.469	1.61e-10	***
x11	1.2030164	0.0996828	12.068	< 2e-16	***
x12	1.9581366	0.0997579	19.629	< 2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1858 on 71 degrees of freedom

Multiple R-squared: 0.9527, Adjusted R-squared: 0.9447

F-statistic: 119.1 on 12 and 71 DF, p-value: < 2.2e-16

From above, we can see that we get same model.