

Problem 2

Q1

```
In [2]: library(quantmod)
library(fpp)
```

Warning message:

"package 'quantmod' was built under R version 3.6.2"

Loading required package: xts

Warning message:

"package 'xts' was built under R version 3.6.2"

Loading required package: zoo

Attaching package: 'zoo'

The following objects are masked from 'package:base':

as.Date, as.Date.numeric

Loading required package: TTR

Warning message:

"package 'TTR' was built under R version 3.6.2"

Registered S3 method overwritten by 'quantmod':

method	from
as.zoo.data.frame	zoo

Version 0.4-0 included new data defaults. See ?getSymbols.

Warning message:

"package 'fpp' was built under R version 3.6.2"

Loading required package: forecast

Warning message:

"package 'forecast' was built under R version 3.6.2"

Registered S3 methods overwritten by 'forecast':

method	from
fitted.fracdiff	fracdiff
residuals.fracdiff	fracdiff

Loading required package: fma

Warning message:

"package 'fma' was built under R version 3.6.2"

Loading required package: expsmooth

Warning message:

"package 'expsmooth' was built under R version 3.6.2"

Loading required package: lmtest

Loading required package: tseries

Warning message:

"package 'tseries' was built under R version 3.6.2"

```
In [3]: getSymbols("SPY",from='2015-01-01',to='2020-02-22',src="yahoo")
dim(SPY)
```

'getSymbols' currently uses auto.assign=TRUE by default, but will use auto.assign=FALSE in 0.5-0. You will still be able to use 'loadSymbols' to automatically load data. getOption("getSymbols.env") and getOption("getSymbols.auto.assign") will still be checked for alternate defaults.

This message is shown once per session and may be disabled by setting options("getSymbols.warning4.0"=FALSE). See ?getSymbols for details.

Warning message:

''indexClass<-' is deprecated.

Use 'tclass<-' instead.

See help("Deprecated") and help("xts-deprecated")."

'SPY'

1293 · 6

From the dimension, we can see that this time series have 1293 days and later we will separate it into a training set (contains 1193 observations) and a test set (contains 100 observations).

```
In [3]: head(SPY)
```

	SPY.Open	SPY.High	SPY.Low	SPY.Close	SPY.Volume	SPY.Adjusted
2015-01-02	206.38	206.88	204.18	205.43	121465900	186.1593
2015-01-05	204.17	204.37	201.35	201.72	169632600	182.7974
2015-01-06	202.09	202.72	198.86	199.82	209151400	181.0756
2015-01-07	201.42	202.72	200.88	202.31	125346700	183.3320
2015-01-08	204.01	206.16	203.99	205.90	147217800	186.5852
2015-01-09	206.40	206.42	203.51	204.25	158567300	185.0900

```
In [4]: tail(SPY)
```

	SPY.Open	SPY.High	SPY.Low	SPY.Close	SPY.Volume	SPY.Adjusted
2020-02-13	335.86	338.12	335.56	337.06	54501900	337.06
2020-02-14	337.51	337.73	336.20	337.60	64582200	337.60
2020-02-18	336.51	337.67	335.21	336.73	57226200	336.73
2020-02-19	337.79	339.08	337.48	338.34	48814700	338.34
2020-02-20	337.74	338.64	333.68	336.95	74163400	336.95
2020-02-21	335.47	335.81	332.58	333.48	113788200	333.48

```
In [5]: data=SPY[,6]
```

```
In [6]: logrtn=dailyReturn(data,type="log")
simplertn=dailyReturn(data)
```

```
In [7]: adf.test(logrtn)
```


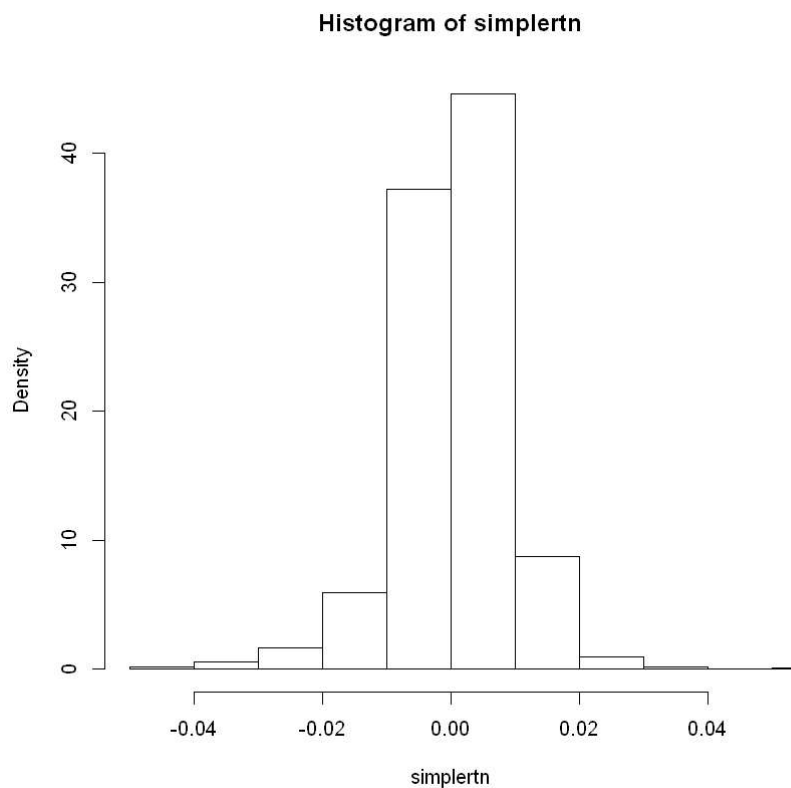
```
Warning message in adf.test(logrtn):  
"p-value smaller than printed p-value"
```

```
Augmented Dickey-Fuller Test
```

```
data: logrtn  
Dickey-Fuller = -11.731, Lag order = 10, p-value = 0.01  
alternative hypothesis: stationary
```

We can see from the outcome of adftest that the $p\text{-value}=0.01<0.05$, so the logrtn time series is stationary at the confidence 95%.

```
In [8]: hist(simplrtn, freq=F, xlim=c(-0.05, 0.05))  
hist(logrtn, freq=F, xlim=c(-0.05, 0.05))  
par(mfrow=c(2, 1))
```



For this time series, we use log to make it stationary.

```
In [9]: data1=logrtn[1:1193]
auto.arima(data1)
```

```
Series: data1
ARIMA(1,0,1) with non-zero mean
```

```
Coefficients:
```

	ar1	ma1	mean
	0.9622	-0.9832	4e-04
s.e.	0.0203	0.0135	1e-04

```
sigma^2 estimated as 7.346e-05: log likelihood=3986.57
AIC=-7965.14 AICc=-7965.11 BIC=-7944.81
```

```
In [10]: m1=arima(data1,order=c(1,0,1))
```

From the auto.arima, we get that our model is arima(1,0,1). This is m1.

```
In [47]: count1=0
for (i in 1:100)
{
data=logrtn[i:1192+i,]
model=arima(data,order=c(1,0,1))
pre=predict(model,1)[1]
if(as.numeric(pre)*(as.numeric(logrtn[1193+i,1]))>0){
  count1=count1+1
}
}
```

[illegible]


```
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
Warning message in arima(data, order = c(1, 0, 1)):
"possible convergence problem: optim gave code = 1"
```

```
In [48]: count1
```

```
49
```

```
In [49]: acc1=count1/100
acc1
```


```
0.49
```

m1=arima(1,0,1) has an accuracy of 0.49.

Q2

MA model

```
In [14]: Acf(data1,lag=15)
```



From the ACF plot, we get two significant order, one is 4, the other is 14. We will compare two models with AIC and error criterion.

```
In [15]: Box.test(logrtn,lag=10)
Box.test(logrtn,lag=10,type="Ljung")
```

Box-Pierce test

```
data: logrtn
X-squared = 14.892, df = 10, p-value = 0.1361
```

Box-Ljung test

```
data: logrtn
X-squared = 14.975, df = 10, p-value = 0.133
```

```
In [16]: m2=arima(data1,c(0,0,4))
```

In [17]: m2

Call:

```
arima(x = data1, order = c(0, 0, 4))
```

Coefficients:

	ma1	ma2	ma3	ma4	intercept
	-0.0247	-0.0531	0.0216	-0.0734	4e-04
s.e.	0.0289	0.0291	0.0295	0.0310	2e-04

sigma^2 estimated as 7.312e-05: log likelihood = 3987.87, aic = -7963.73

In [18]: `Box.test(m2$residuals,lag=10,type="Ljung")`

Box-Ljung test


data: m2\$residuals

X-squared = 7.35, df = 10, p-value = 0.692

In [19]: `res2=m2$resid`
`res2=res2[!is.na(res2)]`
`sd(res2)`

0.00855488918440655

In [20]: `plot(m2$resid)`



In [21]: `m2test=arima(data1,c(0,0,14))`

In [22]: m2test

Call:

```
arima(x = data1, order = c(0, 0, 14))
```

Coefficients:


	ma1	ma2	ma3	ma4	ma5	ma6	ma7	ma8
	-0.028	-0.0454	0.0275	-0.0675	-0.0320	-0.0080	0.0519	-0.0512
s.e.	0.029	0.0291	0.0294	0.0297	0.0294	0.0297	0.0297	0.0284
	ma9	ma10	ma11	ma12	ma13	ma14	intercept	
	-0.0094	-0.0263	-0.0042	0.0203	0.0010	-0.0718	4e-04	
s.e.	0.0288	0.0298	0.0311	0.0293	0.0296	0.0305	2e-04	

sigma^2 estimated as 7.229e-05: log likelihood = 3994.61, aic = -7957.21

```
In [23]: res2test=m2test$resid
res2test=res2test[!is.na(res2test)]
sd(res2test)
```

0.0085062038048902

```
In [24]: plot(m2test$resid)
```



From the AIC and standard variance, we can see that MA(4) is better than MA(14). So we choose $q=4$ and $m2=MA(4)$.

```
In [26]: count2=0
for (i in 1:100)
{
  data=logrtn[i:1192+i,]
  model=arima(data,order=c(0,0,4))
  pre=predict(model,1)[1]
  if(as.numeric(pre)*(as.numeric(logrtn[1193+i,1]))>0){
    count2=count2+1
  }
}
```

```
In [27]: count2

54
```


```
In [28]: acc2=count2/100
acc2

0.54
```

m2=MA(4) has an accuracy of 0.54.

AR model

```
In [29]: pacf(data1, lag=15)
```



From the pacf, we can see that order=4 might be the order of AR. Besides, we will use other functions to help us decide the best order.

```
In [30]: m3=ar(as.vector(data1), method="mle")
```

```
In [31]: m3
```

Call:

```
ar(x = as.vector(data1), method = "mle")
```

Coefficients:

1	2	3	4
-0.0218	-0.0526	0.0152	-0.0626

Order selected 4 sigma^2 estimated as 7.32e-05

The ar function will automatically fit an AR model using AIC criterion. And the order it selected is 4.

```
In [32]: m3=arima(data1,c(4,0,0))
m3
```

Call:

```
arima(x = data1, order = c(4, 0, 0))
```

Coefficients:

	ar1	ar2	ar3	ar4	intercept
	-0.0218	-0.0526	0.0151	-0.0627	4e-04
s.e.	0.0289	0.0289	0.0289	0.0290	2e-04

sigma^2 estimated as 7.319e-05: log likelihood = 3987.3, aic = -7962.61

```
In [35]: count3=0
for (i in 1:100)
{
data=logrtn[i:1192+i,]
model=arima(data,order=c(4,0,0))
pre=predict(model,1)[1]
if(as.numeric(pre)*(as.numeric(logrtn[1193+i,1]))>0){
count3=count3+1
}
}
```

```
In [36]: acc3=count3/100
acc3
```

0.55

As we can see, acc1=0.49, acc2=0.54, acc3=0.55, so we choose m3=AR(4) as our final model.

Q3

```
In [37]: data=logrtn[1:1193,]
model=m3
pre=predict(model,1)[1]
```

```
In [38]: pre
```

\$pred = A Time Series:
0.00124209852240935

```
In [39]: s=100000
for (i in 1:100)
{
data=logrtn[i:1192+i,]

model=arima(data,order=c(4,0,0))

pre=predict(model,1)[1]

if(as.numeric(pre)>0){
  s=s*(1+as.numeric(simplertn[1193+i]))
}
}
```

```
In [40]: s

106803.970424357
```

Q4

```
In [41]: mean=mean(logrtn[1:1193])
mean

0.000382919523466322
```

```
In [42]: s1=100000
for (i in 1:100)
{
data=logrtn[i:1192+i,]

model=arima(data,order=c(4,0,0))

pre=predict(model,1)[1]

if(as.numeric(pre)>mean){
  s1=s1*(1+as.numeric(simplertn[1193+i]))
}
}
```

```
In [43]: s1

104927.815377925
```

Q5


```
In [44]: mean1=mean(logrtn[1:1193])+0.25*sd(logrtn[1:1193])
mean1
0.00253039405171044
```

```
In [45]: s2=100000
for (i in 1:100)
{
data=logrtn[i:1192+i,]
model=arima(data,order=c(4,0,0))
pre=predict(model,1)[1]
if(as.numeric(pre)>mean1){
s2=s2*(1+as.numeric(simplertn[1193+i]))
}
}
```

```
In [46]: s2
1e+05
```

Q6

```
In [16]: try1=0.75*mean(logrtn[1:1193])
try1
0.000287189642599741
```

```
In [44]: n=1
i=0
s=vector()
while(i<1){
s3=100000
try=i*mean(logrtn[1:1193])
for (j in 1:100){
data=logrtn[j:1192+j,]
model=arima(data,order=c(4,0,0))
pre=predict(model,1)[1]
if(as.numeric(pre)>try){
s3=s3*(1+as.numeric(simplertn[1193+j]))
}
}
s=c(s,s3)
n=n+1
i=i+0.1
}
```

In [45]:

s

106803.970424357 · 106803.970424357 · 106150.873439916 · 104397.198249337 ·
 103838.837710216 · 106244.553473954 · 105915.323519614 · 105431.207545448 ·
 104749.539121005 · 105420.231477863 · 104927.815377925

Firstly, we try to write a loop to find out at what level of mean we can get the best return. And we can get above output. As we can see, when $pre > 0$ or $pre > 0.1 * \text{mean}$, the returns are best. So this might be the strategy to find one criterion related to mean. The reason why we don't contain sd is that the standard variance is not at the same quantitative level with mean.

In [28]:

```
s4=100000
for (i in 1:100)
{
  data=logrtn[(892+i):(1192+i),]
  model=arima(data,order=c(4,0,0))
  pre=predict(model,1)[1]
  if(as.numeric(pre)>0){
    s4=s4*(1+as.numeric(simplertn[1193+i]))
  }
}
```

In [29]:

s4

106762.247409162

Secondly, we use the past 300 days to predict and when the pre is bigger than 0, we long. From the output we can see that our return is good. The rational is that the return might more related to recent data so this might be better for our forecast.